



qwen3.8-27b → qwen38_19b_h5120_l48_moe: same hidden size (5120), different computational topology


a Stored parameters



b Active per token



c Depth by block type



d FFN width and MoE sparsity

