

Estimated deployment memory at Q4, fully GPU-resident

ANALYTICAL — computed from qwen_distill.research.memory. No inference memory has been measured on hardware.

